

Task 5: Mmap vs. Pipes vs. Threads

Jana Saleh 900204192
Mariam Dahab 900192441
Muhammad Azzazy 900202821

TABLE OF CONTENTS

01

ROLES

Role of every team member

02

I/O OF MMAP

Detailed description of the inputs and outputs of mmap

03

PSEUDOCODE

Pseudocode of the four compute functions

04

EXAMPLES

Simple examples of the four compute functions producing correct results

05

FIGURES

Figures illustrating the processes, threads, pipes, and the flow of data

06

TESTS DESIGN

How the two tests were designed and the ranges chosen

TABLE OF CONTENTS

07

SETTING N_PROC

How and why n_proc was set

08

SMOOTHING

Method of curve smoothing

09

RESULTS

Comparison results and answers to the two questions

10

Bonus

Comparison results

The background features a dark blue field with glowing cyan circuit-like lines. These lines are composed of straight segments and right-angle turns, with small circular nodes at various points. Some lines run vertically along the left and right edges, while others branch out horizontally or diagonally. The overall aesthetic is high-tech and digital.

01

ROLES

Role of every team member

Name	Role
Jana Saleh	Implemented threads_compute
Mariam Dahab	Plotted the graphs comparing between the four compute functions and implemented the second part of the bonus
Muhammad Azzazy	Implemented mmap_compute



02

I/O OF MMAP



Inputs and outputs of mmap

INPUTS OF MMAP



The inputs of mmap are in the following order:

- The address is set to 0 to allow the system to map anywhere it chooses.
- The size is in our case $(N+2) \times \text{unsigned long int}$, where N is the number of items in the file.
- Some flags such as the following:
 - PROT_READ indicates whether the pages can be read from.
 - PROT_WRITE indicates whether the pages can be written to.
 - MAP_SHARED indicates whether the map is shared with the different processes.
 - MAP_ANONYMOUS indicates that the mapping is not backed by a file and the file descriptor is ignored.
- The file descriptor is set 0 for standard input.
- The offset which must be an integer multiple of the page size is set to 0, but since there is no file it does not matter.

OUTPUTS OF MMAP



The outputs of mmap are in the following order:

- The numbers stored within the file
- The results
- The number of operations done



03

PSEUDOCODE



Pseudocode of the four
compute functions

PSEUDOCODE OF MMAP_COMPUTE

```
MMAP_COMPUTE(n_proc, filepath[], f)
    result ← 0
    Number* head ← NULL
    FILE* fh ← fopen(filepath, "r")
    N ← 0
    while (!feof(fh))
        do
            fscanf(fh, "%lu", &val)
            append(&head, val)
            N ← N + 1

    unsigned long int *ptr ← mmap(NULL, (N+n_proc+1)*sizeof(unsigned long int), PROT_READ | PROT_WRITE, MAP_SHARED |
MAP_ANONYMOUS, 0, 0)
    if ptr == MAP_FAILED
        then printf("Mapping Failed\n")
    for i ← 1 to N
        do
            ptr[i] ← head->num
            head ← head->next
    for i ← 1 to n_proc-1
        do
            pid_t child_pid ← fork()
```

PSEUDOCODE OF MMAP_COMPUTE

```
if child_pid == 0
  then If (N-1)/n_proc == 0
    then ptr[N+i+1] ← (*f)(ptr[N+i], ptr[N+i+1])
    else
      for j ← 1 to (N-1)/n_proc
        do
          ptr[N+i+j+1] ← (*f)(ptr[N+i+j], ptr[N+i+j+1])
          ptr[N+n_proc+1] ← ptr[N+n_proc+1]+1
      exit(0)

wait(NULL)
for k ← ptr[N+n_proc] to N-1
  do
    ptr[N+n_proc] ← ptr[N+n_proc] + 1
    ptr[k+1] ← (*f)(ptr[k], ptr[k+1])
result ← ptr[N]
munmap(ptr, (N+1)*(sizeof(unsigned long int)))
return result
```

PSEUDOCODE OF THREADS_COMPUTE


```
threads_compute( filepath, threads_number, (*f)(int, int)) {  
    file = fopen(filepath, "r")  
    while (fscanf(file, "%d", &m) != EOF)  
        nums[i] = m  
        i++  
    nums = realloc(nums, (i + 1) * sizeof(int))  
    fclose(file)  
    N = i  
    numbers_per_thread = N / threads_number  
    for i ← 0 to threads_number  
        first_num_inthread = i * numbers_per_thread  
        partialsum = 0  
        if i == threads_number - 1  
            last_num_inthread = N  
        else  
            last_num_inthread = (i + 1) * numbers_per_thread
```

```
    for int i ← 0 to threads_number  
        for int j = first_num_inthread to th[i].last  
            numbers_in_thread[count] = numbers[j]  
            count++  
  
    // create threads  
    for i ← 0 to threads_number  
        pthread_create(&threads[i], NULL, threadfun, (void  
        *)&th[i])  
    // wait for threads to terminate  
    for i ← 0 to n_threads  
        pthread_join(threads[i], NULL)  
        pthread_mutex_destroy(&mut);  
    // add the sums of each thread and store it in  
    sum_threads  
    for i ← 0 to n_threads)  
        total_ofThreads = f(total_ofThreads, partialsum)  
    return total_ofThreads
```

A series of glowing cyan lines on the left side of the slide, resembling a circuit board. The lines are vertical and horizontal, with several small cyan dots at the junctions and endpoints. Some lines end in small cyan squares.

04

EXAMPLES

A series of glowing cyan lines on the right side of the slide, resembling a circuit board. The lines are vertical and horizontal, with several small cyan dots at the junctions and endpoints. Some lines end in small cyan squares.A horizontal glowing cyan line spanning the width of the text area, with small cyan dots at each end.

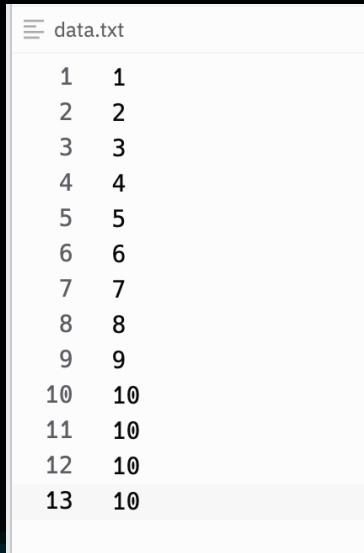
Simple examples of the four
compute functions producing
correct results

MMAP_COMPUTE EXAMPLE

```
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ cc main.c
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 1 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 2 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 3 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 4 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 5 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 6 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 7 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 8 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 9 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 10 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 11 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 12 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 13 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 14 numbers.txt add
15
muhammad@muhammad-OMEN-Laptop-15-en0xxx:~/Desktop$ ./a.out 15 numbers.txt add
15
```

1	1
2	2
3	3
4	4
5	5

EXAMPLE OF THREAD_COMPUTE




data.txt	
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	10
11	10
12	10
13	10

```
./main  
result is 85  
thread Computation took 0.013991 seconds to execute  
>
```

A series of glowing cyan lines on the left side of the slide, resembling a circuit board. The lines are vertical and horizontal, with several small cyan dots at the junctions and endpoints.

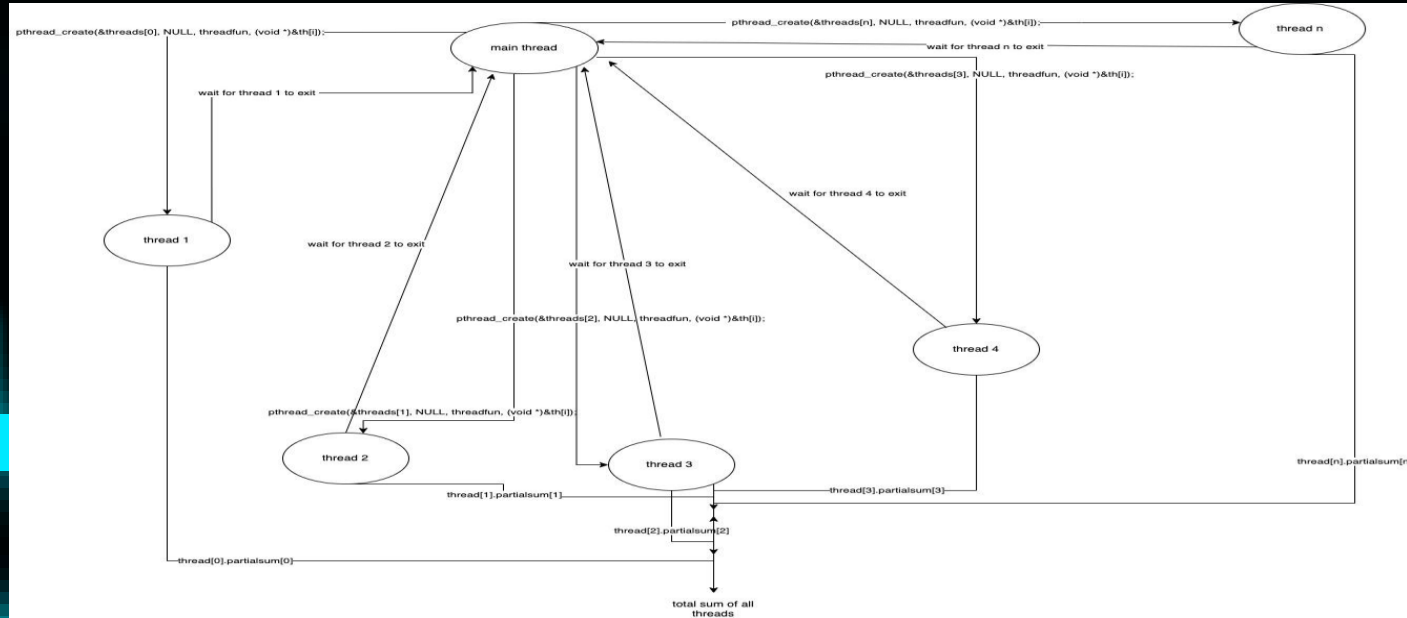
05

FIGURES

A series of glowing cyan lines on the right side of the slide, resembling a circuit board. The lines are vertical and horizontal, with several small cyan dots at the junctions and endpoints.A horizontal glowing cyan line with small cyan dots at each end, positioned above the text.

Figures illustrating the processes, threads, pipes, and the flow of data

Figure of threads dataflow





06

TESTS DESIGN



How the two tests were designed and the ranges chosen

- In the case of fixing N . The range of n_procs was chosen based on the capabilities of the computer (the number of cores) . All functions except sequential where given the same range of n_procs and all of them were give the same file containing N data.
- In case of fixing n_procs . The range of N was set from the least value, which is 10, to the maximum value the computer can handle, which is 10,000. Again in each run, all functions where given the same file containing the same set of values from consistency.
- In both cases the time of execution was gathered 10 times for each combination of numbers are the average was used as the final answer.

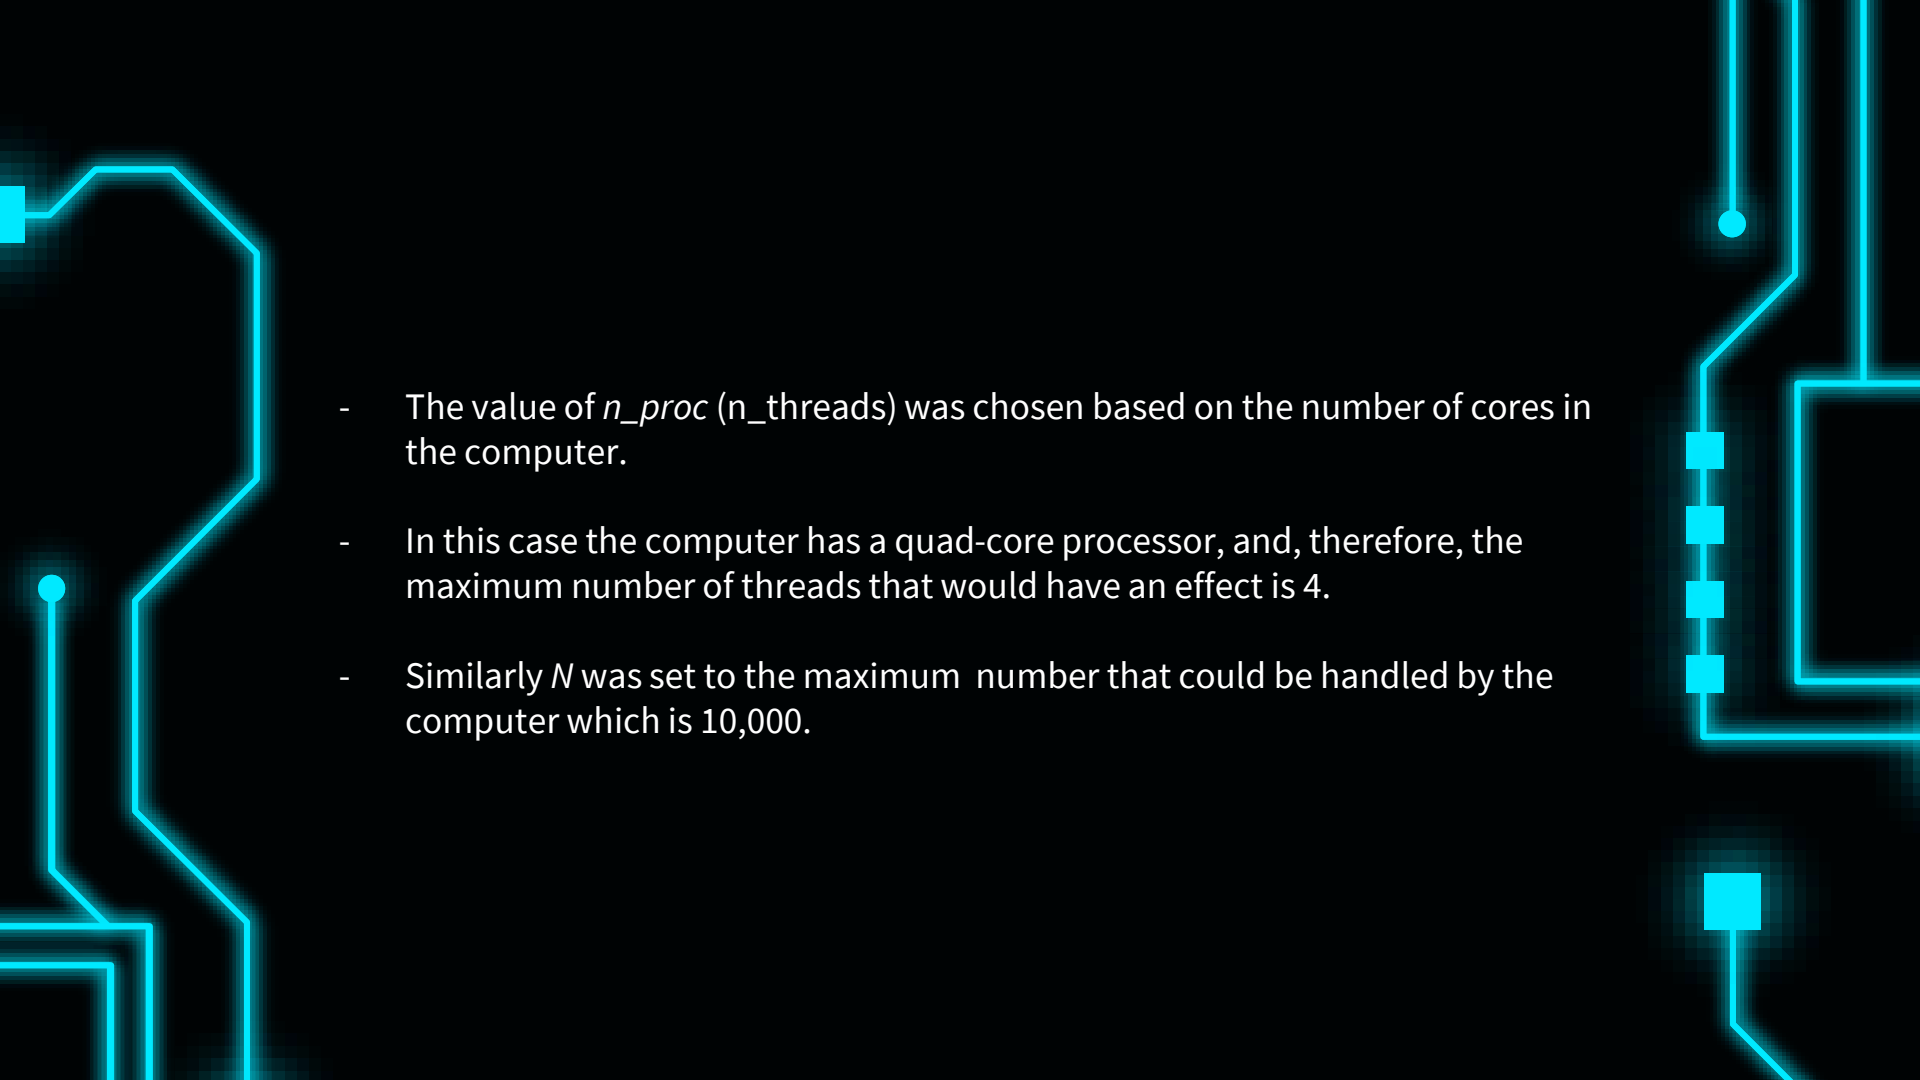


07

FIXING VALUES




Why was each value chosen

- 
- The value of n_{proc} ($n_{threads}$) was chosen based on the number of cores in the computer.
 - In this case the computer has a quad-core processor, and, therefore, the maximum number of threads that would have an effect is 4.
 - Similarly N was set to the maximum number that could be handled by the computer which is 10,000.



08

SMOOTHING




By running each set of values 10 times and taking the average . We were able to create a smooth curve for each line in the curve.



09

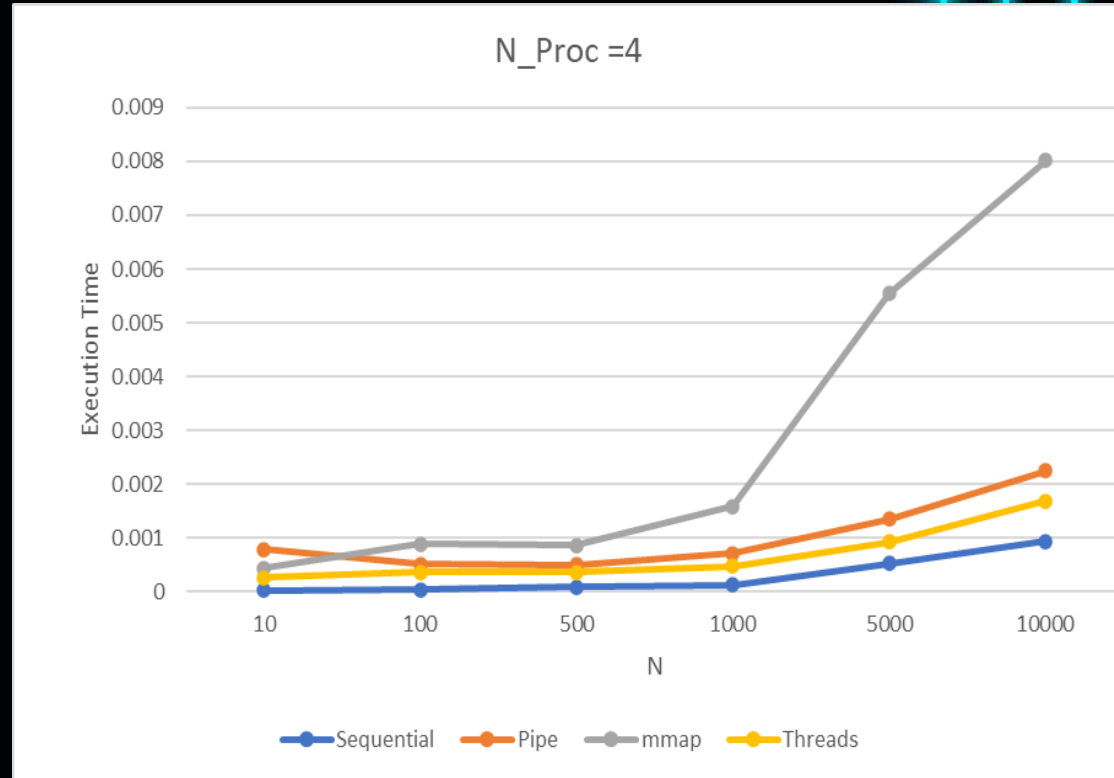
RESULTS



Comparing functions using the graphs

Fixing n_proc / n_threads to 4

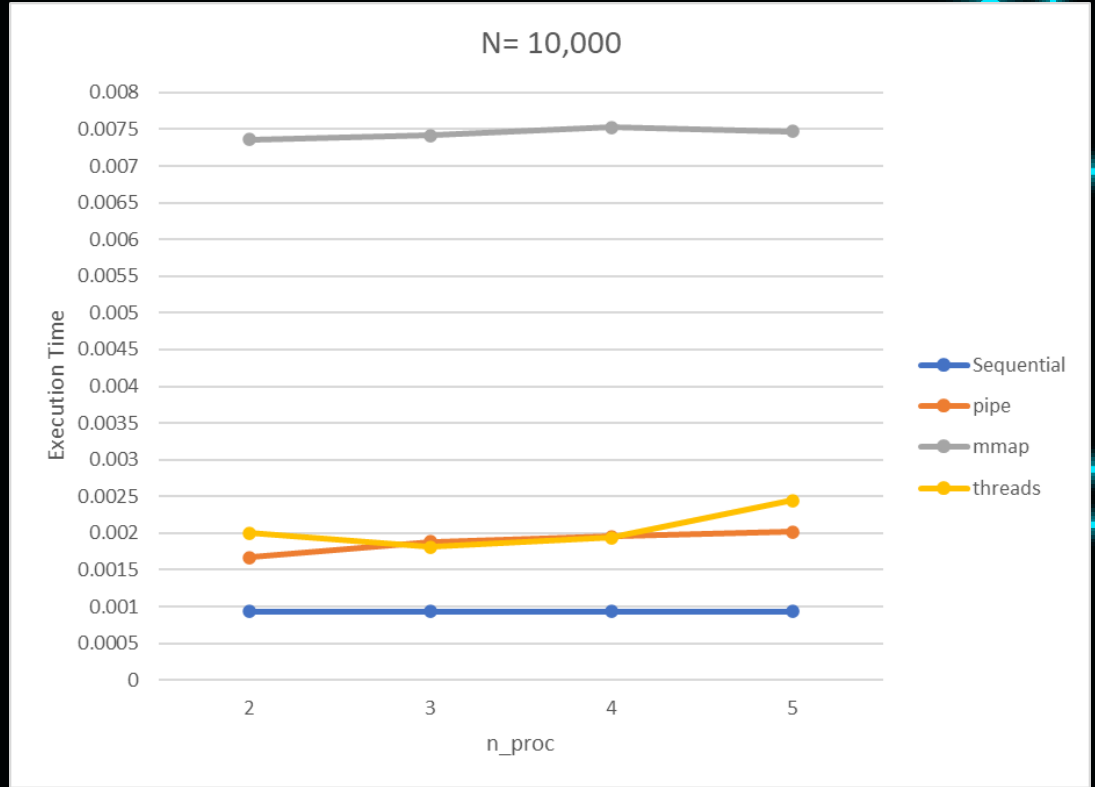
- Although we expected that one of the pipe/thread functions would be the fastest since the data is divided and processed concurrently.
- However, the results different and the sequential took the least time out of the 4 from the beginning.
- Again this is due to the fact that the number of cores limited the dividing of the data and we still believe that using different devices we may get different result.



Fixing N to 10,000

Although we expected that one of the pipe/thread functions would be the fastest since the data is divided and processed concurrently.


- However, similar to the previous graph, the results came different and the sequential took the least time out of the 4 from the beginning.





10


BONUS




Comparing functions using the graphs

Since the single pipe function makes the parent until the created child is done before creating another one it is obvious that in both cases (fixing N and n_proc) the multiple pipe function would outperform since the single can't be considered as parallel computation

N= 10,000



n_proc =4



THANKS!



janasaleh@aucegypt.edu
mhdahab@aucegypt.edu
muhammad-azzazy@aucegypt.edu

Credits: This presentation template was created by
Slidesgo, including icons by **Flaticon**, and
infographics & images by **Freepik**

Please keep this slide for attribution